

# CESC 327

## Assignment 4: Map-Reduce

Professor Oscar Morales

Due date May 11th

### 1 Objective

The objectives of this project are:

- 1) Identify the fundamental concepts of Map Reduce.
- 2) Understand how distributed sorting reduces the sorting time

The task is to develop a map-reduce framework over the distributed file system of the previous lab as we saw in class. It accepts a distributed file where each row has a long and a string separated by comas (Long,String). To store and sort the values in the map and reduce phase use the data structure *TreeMap*. Let  $BMap = TreeMap < Long, List < String > >$  and  $BReduceTreeMap < Long, String >$ .

The interface Context will be the main responsible for checking when the phases have been completed.

```
public interface ContextInterface extends Remote {
    public void setWorkingPeer(Long page);
    public void completePeer(Long page, Long n) throws RemoteException;
    public Boolean isPhaseCompleted();

    public void reduceContext(Long source, ReduceInterface reducer,
        Context context) throws RemoteException;
    public void mapContext(Long page, MapReduceInterface mapper,
        Context context) throws RemoteException;
}
```

```

public interface MapReduceInterface {
    public void map(Long key, String value) throws IOException;
    public void reduce(Long key, String value[]) throws IOException;
}

```

## 2 Remote Procedures

- `emitMap(Long key, String value)`: During the map phase, the implementation of `map` will call `emitMap` in the local peer.

Store `(key,value)` in the peer the responsible with `locateSuccessor(key).emitMap(key, value)`. `emitMap` stores `(key, value)` locally in `BMap`. The implementation of `BTree` does not directly support duplicate keys, but it is possible to handle duplicates by inlining or referencing an object collection as a value.

- `emitReduce(Long key, String value)`: During the map phase, the implementation of `map` will call `emitReduce` in the local peer. No key repetition.

Similar as before, store `(key,value)` in the peer the responsible with `locateSuccessor(key).emitReduce(key, value)`. `emitReduce` stores `(key, value)` locally in `BReduce`.

- `mapContext(Long page, MapReduceInterface mapper, Context context)`: Opens the page (`page`), read line-by-line and execute `mapper.map(key, value, context)`. When it has read the complete file, it calls `context.completePeer(page, n)` where `n` is the number of rows. You have to create a new thread to avoid blocking. Observe that `context` is the instance of the coordinator or initiator.

- `reduceContext(Long source, ReduceInterface reducer, Context context)`: If `source  $\neq$  guid`, call `context.add(guid)` and then call `successor.reduceContext(source, reducer, context)`. Then create a new thread to avoid blocking in which you have to read in order `BReduce`, and execute `reducer.reduce(key, value[], context)`. Make sure that the tree is stored in persistent memory. When it completes reading `BReduce`, it calls `context.complete(guid, n)` where `n` is the number of rows and `guid` is the guid of the peer.

Note: It must exist a metafile called "fileName\_reduce" where fileName is the original logical file that you are sorting with  $n$  pages. Each peer create a page (id) with the data in *BReduce* and insert into "fileName\_reduce".

The implementation of context should look like:

```
public class Context extends ContextInterface {
    Long n = 0;
    Set<Long> set;
    public void setWorkingPeer(Long page)
    {
        set.add(page);
    }
    public void completePeer(Long page, Long n) throws RemoteException;
    {
        this.n += n;
        set.remove(page);
    }
    public Boolean isPhaseCompleted()
    {
        if (set.isEmpty())
            return true;
        return false;
    }
    public void reduceContext(Long source, ReduceInterface reducer,
        Context context) throws RemoteException
    {
        // TODO
    }
    public void mapContext(Long page, MapReduceInterface mapper,
        Context context) throws RemoteException
    {
        // TODO
    }
}
```

Any peer can initialize the service by calling the method *runMapReduce(file)* in the distributed file system. The algorithm is as follows:

```

context= new Context()
mapreduce = new MapReduceInterface();
// map Phases
for each page in metafile.file
    context.add(page);
    let peer be the process responsible for storing page
    peer.mapContext(page, mapreduce, context)
wait until context.hasCompleted() = true
// reduce phase
reduceContext(guid, mapreduce, context);
wait until context.hasCompleted() = true;

```

The following algorithm implement a MapReduce to count the number of words:

```

public class Mapper extends MapReduceInterface {
    public void map(Long key, String value) throws IOException
    {
        For each word in value
            emit(md5(word), word + ":" + 1);
    }

    public void reduce(Long key, String values[]) throws IOException
    {
        word = values[0].split(":")[0]
        emit(key, word + ":" + len(values));
    }
}

```

### 3 Grading

Criteria	Weight
Documentation of your program	15%
Source code (good modularization, coding style, comments)	15%
Execution	70%

The documentation must be generated using Doxygen or Javadocs.

## 4 Deliverables

Compress the source and documentation in a zip file and upload to beachboard. The zipfile must contain two folders:

- 1.- Src: All the source code to compile it
- 2.- Docs: HTML with the documentation. It has to contain the definitions of the methods with a short description, parameters and output of the method.