

# CESC 327

## Assignment 4: Map-Reduce

Professor Oscar Morales

Due date May 11th

### 1 Objective

The objectives of this project are:

- 1) Identify the fundamental concepts of Map Reduce.
- 2) Understand how distributed sorting reduces the sorting time

The task is to develop a map-reduce framework over the distributed file system of the previous lab as we saw in class. It accepts a distributed file where each row has a long and a string separated by comas (Long,String). To store and sort the values in the map and reduce phase use the data structure *TreeMap*. Let  $BMap = TreeMap < Long, List < String > >$  and  $BReduce = TreeMap < Long, String >$ .

We extend the *ChordMessageInterface* to execute the main functionality of the Map Reduce.

```
public interface ChordMessageInterface extends Remote
{
    ...

    public void setWorkingPeer(Long page) throws IOException;
    public void completePeer(Long page, Long n) throws RemoteException;
    public Boolean isPhaseCompleted() throws IOException;
    public void reduceContext(Long source, MapReduceInterface reducer,
        ChordMessageInterface context) throws RemoteException;
    public void mapContext(Long page, MapReduceInterface mapper,
```

```

        ChordMessageInterface context) throws RemoteException;

    public void emitMap(Long key, String value) throws RemoteException;
    public void emitReduce(Long page, List<String> value) throws RemoteException;

}

```

## 2 Remote Procedures

- *emitMap(Longkey, Stringvalue)*: During the map phase, the implementation of map will call *context.emitMap*.  
if (*isKeyInOpenInterval(key, predecessor.getId(), successor.id())*) then store (key,value) locally in *BMap*. Otherwise call *locateSuccessor(key).emitMap(key, value)*. The implementation of BTree does not directly support duplicate keys, but it is possible to handle duplicates by inlining or referencing an object collection as a value.
- *emitReduce(Longpage, List < String > value)*: During the reduce phase, the implementation of reduce will call *context.emitReduce*. if (*isKeyInOpenInterval(key, predecessor.getId(), successor.id())*) then store (key,value) locally in *BReduce*. Otherwise call *locateSuccessor(key).emitReduce(key, value)*. Make sure that the tree is stored in persistent memory.
- *mapContext(Long page, MapReduceInterface mapper, ChordMessageInterface context)*: call *setWorkingPeer(page)* then opens the page (guid), read line-by-line, parse and execute *mapper.map(key, value, context)*. When it has read the complete file, it calls *context.completePeer(page, n)* where *n* is the number of rows. You have to create a new thread to avoid blocking. Observe that context is the instance of the coordinator or initiator.
- *reduceContext(Long source, MapReduceInterface reducer, ChordMessageInterface context)*: If *source*  $\neq$  *guid*, call *successor.reduceContext(source, reducer, context)*. Then, create a new

thread to avoid blocking in which you have to read in order *BMap*, and execute *reducer.reduce(key, value, context)*.

Note: It must exist a metafile called "fileName.reduce" where fileName is the original logical file that you are sorting with *n* pages. Each peer creates a page (guid) with the data in *BReduce* and insert into "fileName.reduce".

The implementation of Chord.java should look like:

```
public class Chord extends java.rmi.server.UnicastRemoteObject
implements ChordMessageInterface
    ...

    private Long numberOfRecords
    private Set<Long> set;
    private Map< Long, List< String >> BMap;
    private Map< Long, String > BReduce;

    public void setWorkingPeer(Long page) throws IOException
    {
        set.add(page);
    }
    public void completePeer(Long page, Long n) throws RemoteException{
        this.numberOfRecords += n;
        set.remove(page);
    }

    public Boolean isPhaseCompleted() throws IOException
    {
        return set.isEmpty();
    }

    public void reduceContext(Long source, MapReduceInterface reducer,
        ChordMessageInterface context) throws RemoteException
    {
        //TOD
```

```

}

public void mapContext(Long page, MapReduceInterface mapper,
    ChordMessageInterface context) throws RemoteException
{
    //TODO
}

public void emitMap(Long key, String value) throws RemoteException
{
    if (isKeyInOpenInterval(key, predecessor.getId(), successor.getId()))
    {
        // insert in the BMap. Allows repetition
        if (!BMap.containsKey(key))
        {
            List< String > list = new List< String >();
            BMap.put(key, list);
        }
        BMap.get(key).add(value);
    }
    else
    {
        ChordMessageInterface peer = this.locateSuccessor(key);
        peer.emitMap(key, value);
    }
}

public void emitReduce(Long key, String value) throws RemoteException
{
    if (isKeyInOpenInterval(key, predecessor.getId(), successor.getId()))
    {
        // insert in the BReduce
        BReduce(key, value);
    }
    else
    {

```

```

        ChordMessageInterface peer = this.locateSuccessor(key);
        peer.emitReduce(key, value);
    }

}

public interface ChordMessageInterface
{

    public void map(Long key, String value,
        ChordMessageInterface context) throws IOException;
    public void reduce(Long key, List< String > value,
        ChordMessageInterface context) throws IOException
};

```

Any peer can initialize the service by calling the method *runMapReduce(file)* in the distributed file system. The algorithm is as follows:

```

mapreduce = new Mapper();
// map Phases
for each page in metafile.file
    let peer be the process responsible for storing page
    peer.mapContext(page, mapreduce, chord)
    wait until context.hasCompleted() = true
// reduce phase
reduceContext(guid, mapreduce, chord);

```

The following algorithm implement a MapReduce to count the number of words:

```

public class Mapper extends MapReduceInterface {
    public void map(Long key, String value,
        ChordMessageInterface context) throws IOException
    {
        context.emitMap(key, value);
    }

    public void reduce(Long key, List< String > value,

```

```

    ChordMessageInterface context) throws IOException
{
    context.emitReduce(key, word + ":" + value.length());
}
}

```

### 3 Grading

| Criteria  | Weight |
|---|--------|
| Documentation of your program                             | 15%    |
| Source code (good modularization, coding style, comments) | 15%    |
| Execution   | 70%    |

The documentation must be generated using Doxygen or Javadocs.

### 4 Deliverables

Compress the source and documentation in a zip file and upload to beachboard. The zipfile must contain two folders:

- 1.- Src: All the source code to compile it
- 2.- Docs: HTML with the documentation. It has to contain the definitions of the methods with a short description, parameters and output of the method.